

### > 1\_write\_cats-vs-dogs\_images.py

- >> Creates in `input/jpg` the folders `test`, `train`, `val`, `calib`
- >> To be executed only once forever (folder `calib` is needed only for Quantization with DeePhi)
- >> This script is different from the one developed for CIFAR10

### > 2a\_create\_lmdb.py

- >> It creates the LMDB databases `input/lmdb/train_lmdb` and `input/lmdb/valid_lmdb` for the training step
- >> To be executed only once forever
- >> This script is different from the one developed for CIFAR10

### > 2b\_compute\_mean.py

- >> It computes the mean values for the `train_lmdb` database in `input/mean.binaryproto`
- >> To be executed only once forever

### > 3\_read\_lmdb.py

- >> Just to debug the first 2 scripts

### > 4\_training.py

- >> To launch the training process in Caffe, given solver and CNN description prototxt files
- >> To be used for any trial of training

### > 5\_plot\_learning\_curve.py + plot\_training\_log.py

- >> To be launched at the end of the training to plot the learning curves of accuracy and lost (in different ways)

### > 6\_make\_predictions.py

- >> To be launched at the end of the training to measure the prediction accuracy achieved by the CNN just trained. You need to have scikit library installed
- >> This script is different from the one developed for CIFAR10

- > All those scripts are orchestrated in the shell script called `caffe_flow_CNN.sh` (see next page). Note that first 4 scripts (1\_\*, 2a\_\*, 2b\_\*, 3\_\*) are commented in the shell, as you run them only once

# caffe\_flow\_AlexNet.sh script

Execute the commands:

- `cd ~/ML/cats-vs-dogs/caffe`
- `source caffe_flow_AlexNet.sh >2&1 | tee logfile_caffe_flow_AlexNet.txt`

```
#!/bin/sh
HOME_DIR=/home/danieleb

CAFFE_ROOT=$HOME_DIR/caffe_tools/Ristretto
CAFFE_TOOLS_DIR=$CAFFE_ROOT/distribute
#working dir
WORK_DIR=$HOME_DIR/ML/cats-vs-dogs/caffe

NUMIT=12000 # number of iterations
NET=alexnetBNnoLRN
MOD_NUM=2 # model number

: '
# #####
# SCRIPTS 1 2 3 (DATABASE AND MEAN VALUES)
echo "DATABASE: training and validation in LMDB, test in JPG and MEAN values"

# prepare the databases
python $WORK_DIR/code/1_write_cats-vs-dogs_images.py -p $HOME_DIR/ML/cats-vs-dogs/input/jpg

#create LMDB databases -training (20K), validation (4K), test (1K) images - and compute mean values
python $WORK_DIR/code/2a_create_lmdb.py -i $HOME_DIR/ML/cats-vs-dogs/input/jpg/ -o $HOME_DIR/ML/cats-vs-dogs/input/lmdb
python $WORK_DIR/code/2b_compute_mean.py -i lmdb/train_lmdb -o mean.binaryproto -w $WORK_DIR/../input/

# #####
# SCRIPT 4 (SOLVER AND TRAINING AND LEARNING CURVE)
echo "TRAINING. Remember that: <Epoch_index = floor((iteration_index * batch_size) / (# data_samples))>"

python $WORK_DIR/code/4_training.py -s ./models/$NET/m$MOD_NUM/solver_$MOD_NUM\_$_NET.prototxt -l ./models/$NET/m$MOD_NUM/logfile_$MOD_NUM\_$_NET.log -c $CAFFE_ROOT


# print image of CNN architecture
echo "PRINT CNN BLOCK DIAGRAM"
python $CAFFE_TOOLS_DIR/python/draw_net.py $WORK_DIR/models/$NET/m$MOD_NUM/train_val_$MOD_NUM\_$_NET.prototxt $WORK_DIR/models/$NET/m$MOD_NUM/bd_$MOD_NUM\_$_NET.png

# #####
# SCRIPT 5: plot the learning curve
echo "PLOT LEARNING CURVES"
python ./code/5_plot_learning_curve.py $WORK_DIR/models/$NET/m$MOD_NUM/logfile_$MOD_NUM\_$_NET.log $WORK_DIR/models/$NET/m$MOD_NUM/plt_train_val_$MOD_NUM\_$_NET.png

python ./code/plot_training_log.py 6 $WORK_DIR/models/$NET/m$MOD_NUM/plt_trainLoss_$MOD_NUM\_$_NET.png $WORK_DIR/models/$NET/m$MOD_NUM/logfile_$MOD_NUM\_$_NET.log
python ./code/plot_training_log.py 2 $WORK_DIR/models/$NET/m$MOD_NUM/plt_testLoss_$MOD_NUM\_$_NET.png $WORK_DIR/models/$NET/m$MOD_NUM/logfile_$MOD_NUM\_$_NET.log
python ./code/plot_training_log.py 0 $WORK_DIR/models/$NET/m$MOD_NUM/plt_testAccuracy_$MOD_NUM\_$_NET.png $WORK_DIR/models/$NET/m$MOD_NUM/logfile_$MOD_NUM\_$_NET.log
# #####
# SCRIPT 6 (PREDICTION)
echo "COMPUTE PREDICTIONS"
python ./code/6_make_predictions.py -d ./models/$NET/m$MOD_NUM/deploy_$MOD_NUM\_$_NET.prototxt -w ./models/$NET/m$MOD_NUM/snapshot_$MOD_NUM\_$_NET\_$_iter_$NUMIT.caffemodel 2>&1 | tee ./models/$NET/m$MOD_NUM/predictions_$MOD_NUM\_$_NET.txt

-:***- caffe_flow_AlexNet.sh All L44 (Shell-script[sh])
```

# The best model is m2: 94% top -1 accuracy



# Solver training parameters

```
File Edit Options Buffers Tools Help
[Icons] Save Undo [Icons]
net: "/home/root2/ML/cats-vs-dogs/caffe/models/alexnetBNnoLRN/m2/train_val_2_alexnetBNnoLRN.prototxt"

test_iter: 80 # test_iter = validation dataset size / validation batch size

test_interval: 1000

base_lr: 0.001

lr_policy: "step"
gamma: 0.1
stepsize: 2500

display: 50

max_iter: 12000
momentum: 0.9
weight_decay: 0.0005

snapshot: 5000

snapshot_prefix: "/home/root2/ML/cats-vs-dogs/caffe/models/alexnetBNnoLRN/m2/snapshot_2_alexnetBNnoLRN_"

solver_mode: GPU

#type: "SGD"
type: "Adam"

random_seed: 1201

name: "alexnetBNnoLRN m2 (as m3 but less DROP and less BN)"
layer {
  name: "data"
  type: "Data"
  top: "data"
  top: "label"
  include {
    phase: TRAIN
  }
  transform_param {
    mirror: true
    crop_size: 227
    mean_value: 106
    mean_value: 116
    mean_value: 124
  }
  data_param {
    source: "/home/root2/ML/cats-vs-dogs/input/lmdb/train_lmdb"
    batch_size: 256
    backend: LMDB
  }
}
layer {
  name: "data"
  type: "Data"
  top: "data"
  top: "label"
  include {
    phase: TEST
  }
  transform_param {
    mirror: false
    crop_size: 227
    mean_value: 106
    mean_value: 116
    mean_value: 124
  }
  data_param {
    source: "/home/root2/ML/cats-vs-dogs/input/lmdb/valid_lmdb"
    batch_size: 50
    backend: LMDB
  }
}
##### CONV1 => BN1/SC1/RELU1 => MAXPOOL1
layer {
  name: "conv1"
  type: "Convolution"
  bottom: "data"
  top: "conv1"
  param {
    lr_mult: 1
    decay_mult: 1
  }
}
```

• type: "Adam"

• batch\_size: 256

-:--- solver\_2\_alexnetBNnoLRN.prototxt All L1 (Fundamental)

-:\*\*\* train\_val\_2\_alexnetBNnoLRN.prototxt Top L43 (Fundamental)

# alexnetBNnoLRN: block diagram

**Note. DeePhi' DPU does not support RELU before BN, as shown in Figure:**

DNNC will generated an error with a misleading message, therefore in the real model they have been swapped.

The real model follow this structure:

CONV1=>BN1=>SC1=>RELU1=>  
POOL1=>DROP1 =>

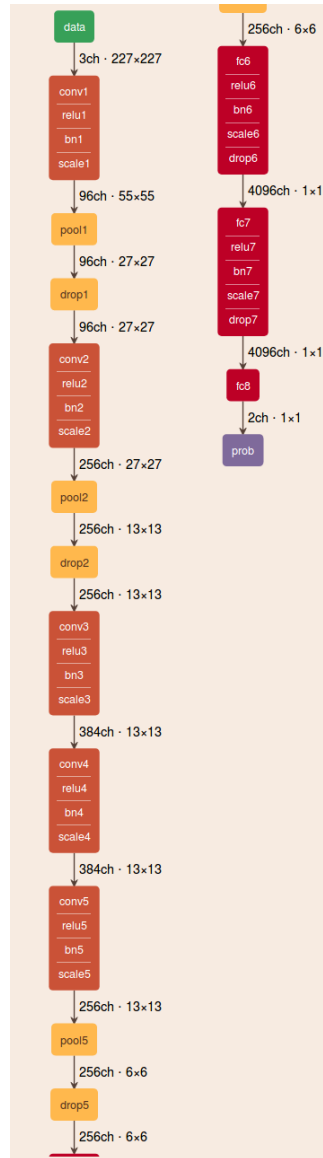
CONV2=>BN2=>SC2=>RELU2=>  
POOL2=>DROP2 =>

CONV3=>BN3=>SC3=>RELU3=>  
CONV4=>BN4=>SC4=>RELU4=>  
CONV5=>BN5=>SC5=>RELU5=>  
POOL5=>DROP5 =>

FC6=>BN6=>SC6=>RELU6=>DROP6 =>

FC7=>BN7=>SC7=>RELU7=>DROP7=>

FC8 =>SOFTMAX



ID	name	type	batch	ch_in	dim_in	ch_out	dim_out	ops	mem
1	data	Input	3	227x227	3	227x227			activation 154.59k
2	conv1	Convolution	3	227x227	96	55x55		mac 105.42M	activation 290.4k param 34.94k
3	relu1	ReLU	96	55x55	96	55x55		comp 290.4k	activation 290.4k
4	bn1	BatchNorm	96	55x55	96	55x55		add 290.4k div 290.4k	activation 290.4k param 192
5	scale1	Scale	96	55x55	96	55x55		mac 290.4k	activation 290.4k
6	pool1	Pooling	96	55x55	96	27x27		comp 629.86k	activation 69.98k
7	drop1	Dropout	96	27x27	96	27x27		comp 69.98k	activation 69.98k
8	conv2	Convolution	96	27x27	256	27x27		mac 447.9M	activation 186.62k param 614.66k
9	relu2	ReLU	256	27x27	256	27x27		comp 186.62k	activation 186.62k
10	bn2	BatchNorm	256	27x27	256	27x27		add 186.62k div 186.62k	activation 186.62k param 512
11	scale2	Scale	256	27x27	256	27x27		mac 186.62k	activation 186.62k
12	pool2	Pooling	256	27x27	256	13x13		comp 389.38k	activation 43.26k
13	drop2	Dropout	256	13x13	256	13x13		comp 43.26k	activation 43.26k
14	conv3	Convolution	256	13x13	384	13x13		mac 149.52M	activation 64.9k param 885.12k
15	relu3	ReLU	384	13x13	384	13x13		comp 64.9k	activation 64.9k
16	bn3	BatchNorm	384	13x13	384	13x13		add 64.9k div 64.9k	activation 64.9k param 768
17	scale3	Scale	384	13x13	384	13x13		mac 64.9k	activation 64.9k
18	conv4	Convolution	384	13x13	384	13x13		mac 224.28M	activation 64.9k param 1.33M
19	relu4	ReLU	384	13x13	384	13x13		comp 64.9k	activation 64.9k
20	bn4	BatchNorm	384	13x13	384	13x13		add 64.9k div 64.9k	activation 64.9k param 768
21	scale4	Scale	384	13x13	384	13x13		mac 64.9k	activation 64.9k
22	conv5	Convolution	384	13x13	256	13x13		mac 149.52M	activation 43.26k param 884.99k

23	relu5	ReLU	256	13x13	256	13x13		comp 43.26k	activation 43.26k
24	bn5	BatchNorm	256	13x13	256	13x13		add 43.26k div 43.26k	activation 43.26k param 512
25	scale5	Scale	256	13x13	256	13x13		mac 43.26k	activation 43.26k
26	pool5	Pooling	256	13x13	256	6x6		comp 82.94k	activation 9.22k
27	drop5	Dropout	256	6x6	256	6x6		comp 9.22k	activation 9.22k
28	fc6	InnerProduct	256	6x6	4096	1x1		mac 37.75M	activation 4.1k param 37.75M
29	relu6	ReLU	4096	1x1	4096	1x1		comp 4.1k	activation 4.1k
30	bn6	BatchNorm	4096	1x1	4096	1x1		add 4.1k div 4.1k	activation 4.1k param 8.19k
31	scale6	Scale	4096	1x1	4096	1x1		mac 4.1k	activation 4.1k
32	drop6	Dropout	4096	1x1	4096	1x1		comp 4.1k	activation 4.1k
33	fc7	InnerProduct	4096	1x1	4096	1x1		mac 16.78M	activation 4.1k param 16.78M
34	relu7	ReLU	4096	1x1	4096	1x1		comp 4.1k	activation 4.1k
35	bn7	BatchNorm	4096	1x1	4096	1x1		add 4.1k div 4.1k	activation 4.1k param 8.19k
36	scale7	Scale	4096	1x1	4096	1x1		mac 4.1k	activation 4.1k
37	drop7	Dropout	4096	1x1	4096	1x1		comp 4.1k	activation 4.1k
38	fc8	InnerProduct	4096	1x1	2	1x1		mac 8.19k	activation 2 param 8.19k
39	prob	Softmax	2	1x1	2	1x1		add 2 div 2 exp 2	activation 2
TOTAL									mac 1.13G comp 1.89M add 658.27k div 658.27k exp 2 activation 3.04M param 58.31M

# alexnetBNnoLRN: Summary

- > Prediction: 94% top-1 average accuracy on the 2 classes
- > 56.88 M parameters in the CNN model (with about 8 layers, counting only the CONV and FC layers)
- > Caffe Training: 12000 iterations @ 256 batch size (20K training and 4K validation images)
  - >> **Elapsed time on GPU P6000 @24GB: 64min**
  - >> Elapsed time on GPU GTX1080 @ 8GB: 162min *(estimated)*
  - >> **Elapsed time on GPU K80 (AWS) @12GB: 152min**
  - >> **Elapsed time on GPU K1000M (\*) @ 2GB: 263min**
- > *(\*) Note: measured on batch\_size=8 and with SGD instead of ADAM to avoid Out Of Memory. Nevertheless, the CNN training process does not converge, as the batch size is too small*