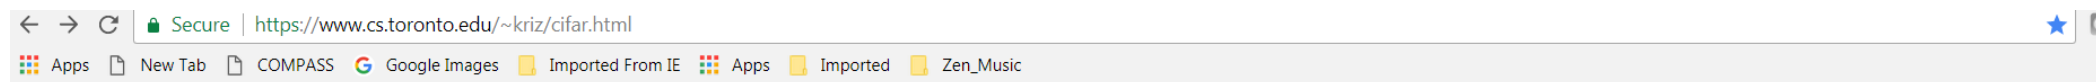


Part 2: Caffe flow with CIFAR10



CIFAR-10 dataset



[< Back to Alex Krizhevsky's home page](#)

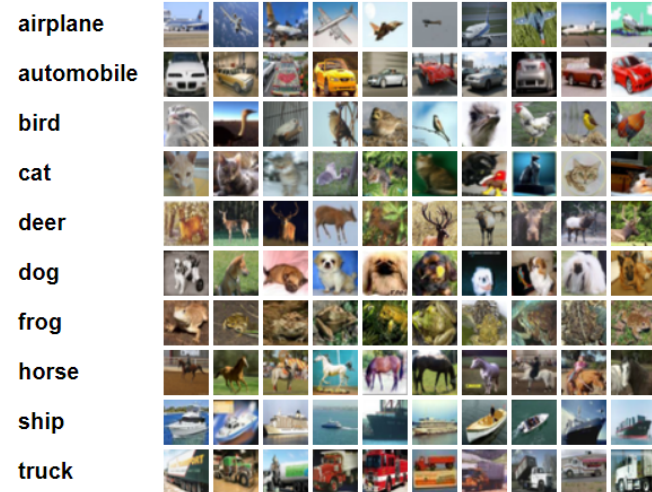
The CIFAR-10 and CIFAR-100 are labeled subsets of the [80 million tiny images](#) dataset. They were collected by Alex Krizhevsky, Vinod Nair, and Geoffrey Hinton.

The CIFAR-10 dataset

The CIFAR-10 dataset consists of 60000 32x32 colour images in 10 classes, with 6000 images per class. There are 50000 training images and 10000 test images.

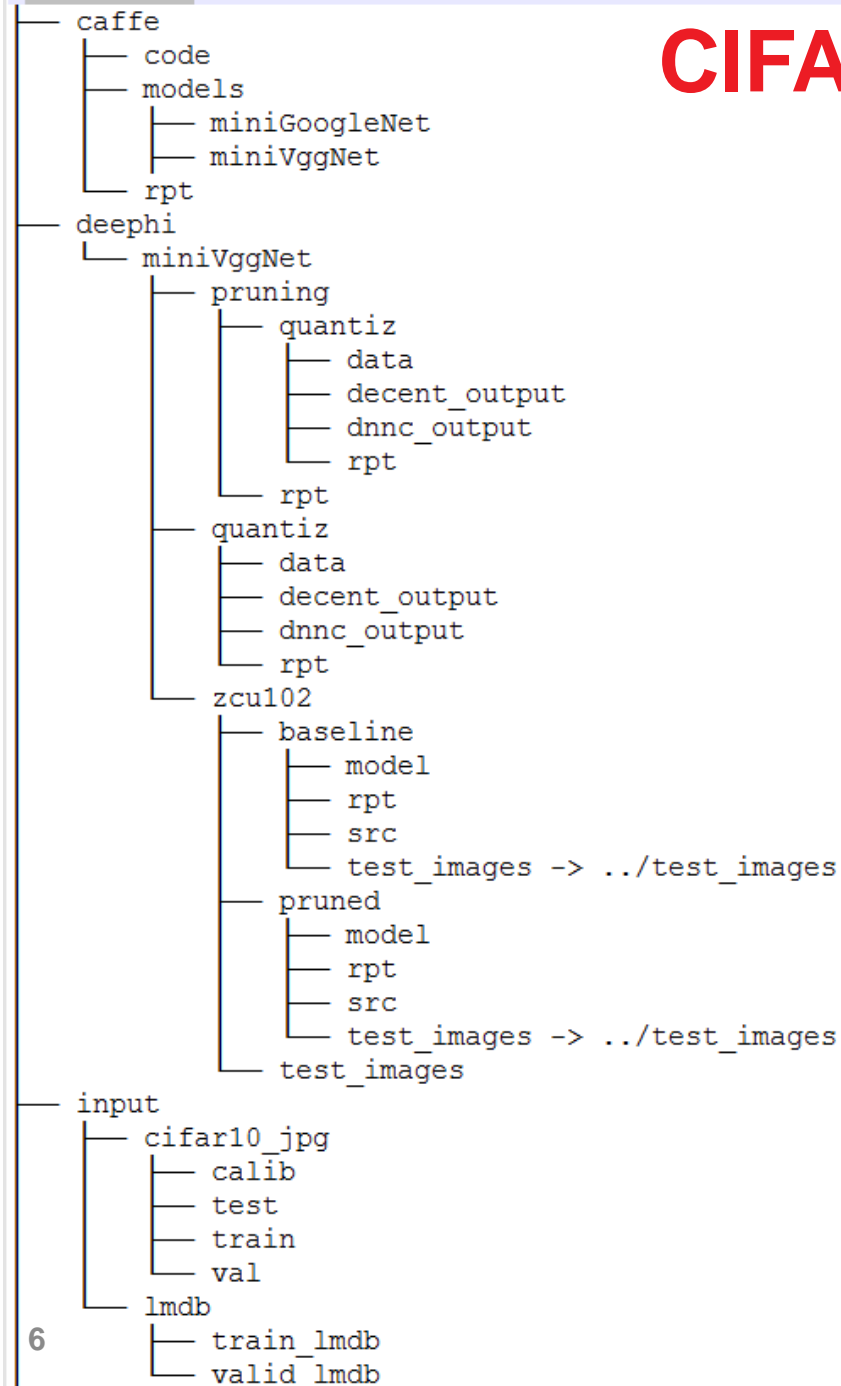
The dataset is divided into five training batches and one test batch, each with 10000 images. The test batch contains exactly 1000 randomly-selected images from each class. The training batches contain the remaining images in random order, but some training batches may contain more images from one class than another. Between them, the training batches contain exactly 5000 images from each class.

Here are the classes in the dataset, as well as 10 random images from each:



The classes are completely mutually exclusive. There is no overlap between automobiles and trucks. "Automobile" includes sedans, SUVs, things of that sort. "Truck" includes only big trucks. Neither includes pickup trucks.

CIFAR-10 project directories structure



The **cifar10** project is organized in the following subdirectories:

- 1) **caffe/code** contains all the Python2.7 scripts
- 2) **caffe/models** contain only the prototxt files for solver, training and deploy of the 2 CNNs
- 3) **caffe/rpt** contains some logfiles as a reference
- 4) **deepphi** contains the files for quantization of baseline CNN (**quantiz**) or of pruned CNN (**pruning**) plus the files for ZCU102 run time execution (respectively **zcu102/baseline**, **zcu102/pruned**)
- 5) **input** contains the JPEG images and lmdb databases for Caffe phases of training and validation and JPEG images for testing top-1 accuracy, plus the JPEG images for DeePhi' calibration during the quantization phase

Python 2.7 scripts

> 1_write_cifar10_images.py

- >> download the dataset from `keras.datasets` (you must have keras/tensorflow installed into your python virtual env, it is needed only to run this script)
- >> store it in JPEG format on the folder `input/cifar10_jpg` (`test`, `train`, `val`, `calib`)
- >> To be executed only once forever (folder `calib` is needed only for Quantization with DeePhi)

> 2a_create_lmdb.py

- >> It creates the LMDB databases `input/lmdb/train_lmdb` and `input/lmdb/valid_lmdb` for the training step
- >> To be executed only once forever

> 2b_compute_mean.py

- >> It computes the mean values for the `train_lmdb` database in `input/mean.binaryproto`
- >> To be executed only once forever

> 3_read_lmdb.py

- >> Just to debug the first 2 scripts

Python 2.7 scripts

- > **4_training.py**
 - >> To launch the training process in Caffe, given solver and CNN description prototxt files
 - >> To be used for any trial of training
- > **5_plot_learning_curve.py + plot_training_log.py**
 - >> To be launched at the end of the training to plot the learning curves of accuracy and lost (in different ways)
- > **6_make_predictions.py**
 - >> To be launched at the end of the training to measure the prediction accuracy achieved by the CNN just trained. You need to have scikit library installed
- > **check_dpu_runtime_accuracy.py**
 - >> To be launched only once the CNN will be running on the ZCU102 board, to compute the effective top-1 accuracy of DeePhi DPU at runtime
- > All those scripts can be orchestrated in the shell script called **caffe_flow_CNN.sh** (see next page), and normally the first 4 scripts are commented there

caffe_flow_CNN.sh script

```
1  #!/bin/sh
2  CAFFE_ROOT=/caffe/BVLC1v0-Caffe
3  CAFFE_TOOLS_DIR=$CAFFE_ROOT/distribute
4  WORK_DIR=/home/ML/cifar10/caffe #working dir
5
6  MOD_NUM=3 # model number
7  NUMIT=40000 # number of iterations
8  NET=miniVggNet
9
10 # #####
11 # SCRIPTS 1 2 3 (DATABASE AND MEAN VALUES)
12 echo "DATABASE: training and validation in LMDB, test in JPG and MEAN values"
13
14 # load the database from keras and write it as JPEG images
15 python $WORK_DIR/code/1_write_cifar10_images.py
16
17 #create LMDB databases -training (50K), validation (9K), test (1K) images - and compute mean values
18 python $WORK_DIR/code/2a_create_lmdb.py
19 python $WORK_DIR/code/2b_compute_mean.py
20
21 #check goodness of LMDB databases (just for debug: you can skip it)
22 python $WORK_DIR/code/3_read_lmdb.py
23
24 # #####
25 # SCRIPT 4 (SOLVER AND TRAINING AND LEARNING CURVE)
26 echo "TRAINING. Remember that: <Epoch_index = floor((iteration_index * batch_size) / (# data_samples))>"
27
28 python $WORK_DIR/code/4_training.py -s ./models/$NET/m$MOD_NUM/solver_$MOD_NUM\_NET.prototxt -l ./models/$NET/m$MOD_NUM/logfile_$MOD_NUM\_NET.log
29
30 # print image of CNN architecture
31 echo "PRINT CNN BLOCK DIAGRAM"
32 python $CAFFE_TOOLS_DIR/python/draw_net.py $WORK_DIR/models/$NET/m$MOD_NUM/train_val_$MOD_NUM\_NET.prototxt $WORK_DIR/models/$NET/m$MOD_NUM/bd_$MOD_NUM\_NET.png
33
34 # #####
35 # SCRIPT 5: plot the learning curve
36 echo "PLOT LEARNING CURVES"
37 python $WORK_DIR/code/5_plot_learning_curve.py $WORK_DIR/models/$NET/m$MOD_NUM/logfile_$MOD_NUM\_NET.log $WORK_DIR/models/$NET/m$MOD_NUM/plt_train_val_$MOD_NUM\_NET.png
38
39 # #####
40 # SCRIPT 6 (PREDICTION)
41 echo "COMPUTE PREDICTIONS"
42 python $WORK_DIR/code/6_make_predictions.py -d $WORK_DIR/models/$NET/m$MOD_NUM/deploy_$MOD_NUM\_NET.prototxt \
43 -w $WORK_DIR/models/$NET/m$MOD_NUM/snapshot_$MOD_NUM\_NET\_iter_$NUMIT.caffemodel 2>&1 | tee $WORK_DIR/models/$NET/m$MOD_NUM/predictions_$MOD_NUM\_NET.txt
44
```

AWS AMI for ML

Select Deep Learning AMI (Ubuntu)

The screenshot shows the AWS Management Console interface for the 'Launch Instance Wizard' in the 'eu-west-1' region. The wizard is at 'Step 1: Choose an Amazon Machine Image (AMI)'. On the left sidebar, under 'Linux (3)', the 'Ubuntu (51)' option is selected. The main area displays a list of AMIs. The 'Deep Learning AMI (Ubuntu)' is highlighted with a red box. It is an Amazon Web Services AMI with a rating of 5 stars (6 reviews), updated on 9/25/18, and costs \$0.0208 to \$41.944/hr including EC2 charges. It includes popular deep learning frameworks like Apache MXNet, TensorFlow, PyTorch, Caffe, and Caffe2. Other visible AMIs include 'Accelerated Machine Learning' by InAccel and 'Predictive Analytics Framework' by RRecktek.

Step 1: Choose an Amazon Machine Image (AMI)

Cancel and Exit

1. Choose AMI 2. Choose Instance Type 3. Configure Instance 4. Add Storage 5. Add Tags 6. Configure Security Group 7. Review

Amazon Linux (75)
Debian (2)
CentOS (11)
Red Hat Enterprise Linux (3)
☒ Ubuntu (51)
Other Linux (2)

Software Pricing Plans
☐ Free (18)
☐ Hourly (29)
☐ Annual (8)
☐ Bring Your Own License (4)

Software Free Trial
☐ Free Trial (16)

Region
☒ Current Region (51)
☐ All Regions (722)

Deep Learning AMI (Ubuntu)
★★★★★ (6) | 15.0 | By Amazon Web Services
\$0.0208 to \$41.944/hr incl EC2 charges + other AWS usage fees
Linux/Unix, Ubuntu 16.04 | 64-bit Amazon Machine Image (AMI) | Updated: 9/25/18
Deep Learning AMI comes with popular deep learning frameworks optimized for high performance on Amazon EC2 instances. Includes Apache MXNet, TensorFlow, PyTorch, Caffe, Caffe2, ...
[More info](#)

Accelerated Machine Learning
★★★★★ (0) | AML_v1.0 | By InAccel
\$2.00/hr for software + AWS usage fees
Linux/Unix, Ubuntu 16.04 | 64-bit Amazon Machine Image (AMI) | Updated: 3/9/18
AML is InAccel's accelerated machine learning library. It aims to maintain the practical and easy to use interface of other open source frameworks, i.e. of Apache Spark, and at the ...
[More info](#)

Predictive Analytics Framework
★★★★★ (4) | 2.3.0 | Previous versions | By RRecktek
\$1.00 to \$2.85/hr for software + AWS usage fees
Linux/Unix, Ubuntu 16.04.03 LTS | 64-bit Amazon Machine Image (AMI) | Updated: 7/17/18
Optimized and complete data analysis environment in R and Python. Specially optimized (OPMI) version of R 3.5.1 with 12583+ installed packages substantially outperforms "free R". ...
[More info](#)

ADAPA Decision Engine
★★★★★ (3) | 5.0.1 | By Zementis
Starting from \$0.99/hr or from \$6,937/yr (20% savings) for software + AWS usage fees
Linux/Unix, Ubuntu 14.04x64 | 64-bit Amazon Machine Image (AMI) | Updated: 6/11/18
ADAPA is a predictive analytics decision engine based on the PMML (Predictive Model Markup Language) standard. With ADAPA, deploy one or many predictive models from data mining ...
[More info](#)

AWS EC2 instance

Launch Instance ▾

Connect

Actions ▾

Filter by tags and attributes or search by keyword

<input type="checkbox"/>	Name ▾	Instance ID ▴	Instance Type ▾	Availability Zone ▾	Instance State ▾	Status Checks ▾	Alarm Status	Public DNS (IPv4) ▾	IPv4 Public IP
		i-0f0652a6b93847cd7	p2.xlarge	eu-west-1b	pending	Initializing	None		-

Instance: **i-0f0652a6b93847cd7** Private IP: 172.31.4.119

Description

Status Checks

Monitoring

Tags

Instance ID

Instance state

Instance type

Elastic IPs

Availability zone

Security groups

Scheduled events

AMI ID

Platform

IAM role

i-0f0652a6b93847cd7

pending

p2.xlarge

eu-west-1b

[Deep Learning AMI -Ubuntu--15-0-AutogenByAWSMP-](#). [view inbound rules](#). [view outbound rules](#)

-

[Deep Learning AMI \(Ubuntu\) Version 15.0 \(ami-028a25f96614a39e9\)](#)

-

-

Public DNS (IPv4)

IPv4 Public IP

IPv6 IPs

Private DNS

Private IPs

Secondary private IPs

VPC ID

Subnet ID

Network interfaces

Source/dest. check

-

-

-

ip-172-31-4-119.eu-west-1.compute.internal

172.31.4.119

vpc-5127a337

subnet-1bdd087d

[eth0](#)

True

Logging in into AWS

```
ec2-user@ip-172-31-10-57:~  
[ec2-user@ip-172-31-10-57 ~]$  
[ec2-user@ip-172-31-10-57 ~]$ cd  
[ec2-user@ip-172-31-10-57 ~]$ ls -l  
total 2800  
drwxrwxr-x 24 ec2-user ec2-user 4096 19 lug 19.19 anaconda3  
drwxrwxr-x 8 ec2-user ec2-user 4096 19 lug 18.10 examples  
drwxrwxr-x 3 ec2-user ec2-user 4096 4 ott 06.40 ML  
-rw-rw-r-- 1 ec2-user ec2-user 2839191 14 lug 02.18 Nvidia_Cloud_EULA.pdf  
-rw-rw-r-- 1 ec2-user ec2-user 2792 12 ago 13.50 README  
drwxrwxr-x 10 ec2-user ec2-user 4096 19 lug 20.02 src  
drwxrwxr-x 5 ec2-user ec2-user 4096 19 lug 18.10 tutorials  
[ec2-user@ip-172-31-10-57 ~]$ ls -l ML  
total 73984  
drwxrwxr-x 5 ec2-user ec2-user 4096 10 set 16.41 cats-vs-dogs  
-rw-rw-r-- 1 ec2-user ec2-user 75751719 4 ott 06.40 ML-CIFAR10-tutorial-master.zip  
[ec2-user@ip-172-31-10-57 ~]$ ls -l src/caffe_python_2/  
total 124  
drwxrwxr-x 14 ec2-user ec2-user 4096 19 lug 20.02 build  
-rw-rw-r-- 1 ec2-user ec2-user 1180 19 lug 20.01 caffe.cloc  
drwxrwxr-x 5 ec2-user ec2-user 4096 19 lug 20.01 cmake  
-rw-rw-r-- 1 ec2-user ec2-user 4197 19 lug 20.01 CMakeLists.txt  
-rw-rw-r-- 1 ec2-user ec2-user 1917 19 lug 20.01 CONTRIBUTING.md  
-rw-rw-r-- 1 ec2-user ec2-user 620 19 lug 20.01 CONTRIBUTORS.md  
drwxrwxr-x 5 ec2-user ec2-user 4096 19 lug 20.01 data  
drwxrwxr-x 4 ec2-user ec2-user 4096 19 lug 20.01 docker  
drwxrwxr-x 6 ec2-user ec2-user 4096 19 lug 20.01 docs  
drwxrwxr-x 15 ec2-user ec2-user 4096 19 lug 20.01 examples  
drwxrwxr-x 3 ec2-user ec2-user 4096 19 lug 20.01 include  
-rw-rw-r-- 1 ec2-user ec2-user 210 19 lug 20.01 INSTALL.md  
-rw-rw-r-- 1 ec2-user ec2-user 2092 19 lug 20.01 LICENSE  
-rw-rw-r-- 1 ec2-user ec2-user 24041 19 lug 20.01 Makefile  
-rw-rw-r-- 1 ec2-user ec2-user 4634 19 lug 20.01 Makefile.config  
-rw-rw-r-- 1 ec2-user ec2-user 4631 19 lug 20.01 Makefile.config.example  
drwxrwxr-x 5 ec2-user ec2-user 4096 19 lug 20.01 matlab  
drwxrwxr-x 7 ec2-user ec2-user 4096 19 lug 20.01 models  
drwxrwxr-x 3 ec2-user ec2-user 4096 19 lug 20.01 python  
-rw-rw-r-- 1 ec2-user ec2-user 2130 19 lug 20.01 README.md  
drwxrwxr-x 3 ec2-user ec2-user 4096 19 lug 20.01 scripts  
drwxrwxr-x 4 ec2-user ec2-user 4096 19 lug 20.01 src  
drwxrwxr-x 3 ec2-user ec2-user 4096 19 lug 20.01 tools
```

It is assumed that:

- 1) you have created a folder named **~/ML**
- 2) You have unzipped the **ML-CIFAR10-tutorial-master.zip**
- 3) You have renamed it simply **cifar10**
- 4) At the end of above steps the folder **~/ML/cifar10** exists

WARNING: Caffe with Python2.7 (caffe_p27) on AWS

- > Caffe on AWS has a slightly different installation structure, being based on **anaconda**. In particular:
- folder **distribute** does not exist and is replaced by folder **build/install**
 - **caffe.bin** executable is just named **caffe** as in previous page
 - **compute_image_mean.bin** has not been compiled with OpenCV: you should recompile the whole Caffe on AWS, do not do that, too much time consuming! Instead, just copy `mean.binaryproto` I have provided in the archive to your AWS site

WARNING: Tensorflow/Python2.7 (tensorflow_p27) on AWS

- > Furthermore, from **tensorflow_p27** environment (source activate tensorflow_p27) I also had to run the following commands (the first one only once forever)

```
>> conda install keras
>> export LD_LIBRARY_PATH=~/.src/caffe_python_2/build/install/lib64/:$LD_LIBRARY_PATH
>> export PYTHONPATH=/usr/local/lib/python2.7/dist-packages/:$PYTHONPATH
```

- > **ALL THOSE SETTINGS ARE CAPTURED IN THE SCRIPT set_aws_ML_env_cuda9.sh** (have a look at it before launching, read the comments!)

```
>> cd ~/ML/cifar10/aws_scripts
>> source set_aws_ML_env_cuda9.sh 2>&1 | tee ../rpt/logfile_set_aws_ML_env.txt
```

aws_caffe_flow_miniVggNet.sh script

```
1  #!/bin/sh
2  CAFFE_ROOT=/caffe/BVLC1v0-Caffe
3  CAFFE_TOOLS_DIR=$CAFFE_ROOT/distribute
4  WORK_DIR=/home/ML/cifar10/caffe #working dir
5  MOD_NUM=3 # model number
6  NUMIT=40000 # number of iterations
7  NET=miniVggNet
8
9  # #####
10 # SCRIPTS 1 2 3 (DATABASE AND MEAN VALUES)
11
12 # go to TensorFlow environment, load the database from Keras and write it as JPEG images
13 source activate tensorflow_p27
14 export PYTHONPATH=/usr/local/lib/python2.7/dist-packages:/home/ubuntu/anaconda3/envs/caffe_p27/lib/python2.7/site-packages/:$PYTHONPATH #needed for opencv
15 echo "DATABASE: training and validation in LMDB, test in JPG and MEAN values"
16 python $WORK_DIR/code/1_write_cifar10_images.py --pathname /home/danieleb/ML/cifar10/input/cifar10_jpg
17 export PYTHONPATH=/home/ubuntu/src/cntk/bindings/python
18 source deactivate tensorflow_p27
19
20 # go to Caffe environment
21 source activate caffe_p27
22
23 #create LMDB databases -training (50K), validation (9K), test (1K) images - and compute mean values
24 python $WORK_DIR/code/2a_create_lmdb.py
25 # DO NOT RUN BELOW COMMAND AS mean.binaryproto is already available! Is is here only for reference
26 ##python $WORK_DIR/code/2b_compute_mean.py
27
28 #check goodness of LMDB databases (just for debug: you can skip it)
29 python $WORK_DIR/code/3_read_lmdb.py
30
31 # #####
32 # SCRIPT 4 (SOLVER AND TRAINING AND LEARNING CURVE)
33 echo "TRAINING. Remember that: <Epoch_index = floor((iteration_index * batch_size) / (# data_samples))>"
34 python $WORK_DIR/code/4_training.py -s ./models/$NET/m$MOD_NUM/solver_$MOD_NUM\_NET.prototxt -l ./models/$NET/m$MOD_NUM/logfile_$MOD_NUM\_NET.log
35
36 # print image of CNN architecture
37 echo "PRINT CNN BLOCK DIAGRAM"
38 python $CAFFE_TOOLS_DIR/python/draw_net.py $WORK_DIR/models/$NET/m$MOD_NUM/train_val_$MOD_NUM\_NET.prototxt $WORK_DIR/models/$NET/m$MOD_NUM/bd_$MOD_NUM\_NET.png
39
40 # #####
41 # SCRIPT 5: plot the learning curve
42 echo "PLOT LEARNING CURVERS"
43 python ./code/5_plot_learning_curve.py $WORK_DIR/models/$NET/m$MOD_NUM/logfile_$MOD_NUM\_NET.log $WORK_DIR/models/$NET/m$MOD_NUM/plt_train_val_$MOD_NUM\_NET.png
44
```