

# Part 3a. Training with Caffe: miniVggNet



# What is miniVggNet

- > It is a CNN inspired by the original VGG, suitable for the CIFAR10 dataset
- > **It is published on the great book “Deep Learning for Computer Vision with Python: Practitioner Bundle” by Dr. Adrian Rosebrock (pyimagesearch), briefly cited as [DL4CV\_PB]**
- > Original description and training scripts are available only in Keras/Tensorflow
- > Prediction accuracy in the range of 80-82%, depending on how it is measured
- > For Caffe, I re-wrote the prototxt description and solver files and trained the CNN from scratch, doing the proper changes as Caffe and Keras are pretty different environments for ML development
- > To get the same results (within 2% more or less) shown in the next pages, run the following script:  

```
>> source caffe_flow_miniVggNet.sh
```

# MiniVggNet model in Keras, from [DL4CV\_PB]

- > MiniVggNet consists of
  - >> two sets of CONV=>RELU=>CONV=>RELU=>POOL
  - >> followed by a set of FC=>RELU=>FC=>SOFTMAX layers
  - >> The first 2 CONV layers will learn 32 filters of 3x3 size
  - >> The second 2 CONV layers will learn 64 filters again of 3x3 size
  - >> POOL performs pooling over a 2x2 window with a 2x2 stride
  - >> Batch Normalization (BN) are inserted after the RELU activations (ACT) while DROPOUT layers after the POOL and FC layers

Layer Type	Output Size	Filter Size / Stride
INPUT IMAGE	$32 \times 32 \times 3$	
CONV	$32 \times 32 \times 32$	$3 \times 3, K = 32$
ACT RELU	$32 \times 32 \times 32$	
BN	$32 \times 32 \times 32$	
CONV	$32 \times 32 \times 32$	$3 \times 3, K = 32$
ACT RELU	$32 \times 32 \times 32$	
BN	$32 \times 32 \times 32$	
POOL	$16 \times 16 \times 32$	$2 \times 2$
DROPOUT	$16 \times 16 \times 32$	
CONV	$16 \times 16 \times 64$	$3 \times 3, K = 64$
ACT RELU	$16 \times 16 \times 64$	
BN	$16 \times 16 \times 64$	
CONV	$16 \times 16 \times 64$	$3 \times 3, K = 64$
ACT RELU	$16 \times 16 \times 64$	
BN	$16 \times 16 \times 64$	
POOL	$8 \times 8 \times 64$	$2 \times 2$
DROPOUT	$8 \times 8 \times 64$	
FC	512	
ACT RELU	512	
BN	512	
DROPOUT	512	
FC	10	
SOFTMAX	10	

# MiniVggNet: Python code from [DL4CV\_PB]

```
# import the necessary packages
from keras.models import Sequential
from keras.layers.normalization import BatchNormalization
from keras.layers.convolutional import Conv2D
from keras.layers.convolutional import MaxPooling2D
from keras.layers.core import Activation
from keras.layers.core import Flatten
from keras.layers.core import Dropout
from keras.layers.core import Dense
from keras import backend as K

class MiniVGGNet:
    @staticmethod
    def build(width, height, depth, classes):
        # initialize the model along with the input shape to be
        # "channels last" and the channels dimension itself
        model = Sequential()
        inputShape = (height, width, depth)
        chanDim = -1

        # if we are using "channels first", update the input shape
        # and channels dimension
        if K.image_data_format() == "channels_first":
            inputShape = (depth, height, width)
            chanDim = 1
```

```
# first CONV => RELU => CONV => RELU => POOL layer set
model.add(Conv2D(32, (3, 3), padding="same",
                 input_shape=inputShape))
model.add(Activation("relu"))
model.add(BatchNormalization(axis=chanDim))
model.add(Conv2D(32, (3, 3), padding="same"))
model.add(Activation("relu"))
model.add(BatchNormalization(axis=chanDim))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

# second CONV => RELU => CONV => RELU => POOL layer set
model.add(Conv2D(64, (3, 3), padding="same"))
model.add(Activation("relu"))
model.add(BatchNormalization(axis=chanDim))
model.add(Conv2D(64, (3, 3), padding="same"))
model.add(Activation("relu"))
model.add(BatchNormalization(axis=chanDim))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))

# first (and only) set of FC => RELU layers
model.add(Flatten())
model.add(Dense(512))
model.add(Activation("relu"))
model.add(BatchNormalization())
model.add(Dropout(0.5))

# softmax classifier
model.add(Dense(classes))
model.add(Activation("softmax"))

# return the constructed n
return model
```

Keras assumes  
implicitly  
stride is equal to  
max pooling size

## 7



# miniVggNet: deploy and train\_val prototxt files

- > Remember that there are differences between **train\_val\_3\_miniVggNet.prototxt** and **deploy\_3\_miniVggNet.prototxt** description files:
  - >> In the training you need to declare the TRAIN and TEST phases and proper databases, while in the deploy you need to remove those layers
  - >> Batch Normalization layers require the following setting:
    - (deploy) `use_global_stats: true`
    - (train\_val) `use_global_stats: false`
  - >> The last layer – Loss – requires the following setting
    - (deploy) `name: "prob" type: "Softmax" top: "prob"`
    - (train\_val) `name: "loss" type: "SoftmaxWithLoss" top: "loss"`

# train\_val\_3\_miniVggNet.prototxt fragments

```
layer {
  name: "data"
  type: "Data"
  top: "data"
  top: "label"
  include {
    phase: TRAIN
  }
  transform_param {
    mirror: true
    mean_file: "/home/ML/cifar10/input/mean.binaryproto"
  }
  data_param {
    source: "/home/ML/cifar10/input/lmdb/train_lmdb"
    batch_size: 128
    backend: LMDB
  }
}
layer {
  name: "data"
  type: "Data"
  top: "data"
  top: "label"
  include {
    phase: TEST
  }
  transform_param {
    mirror: false
    mean_file: "/home/ML/cifar10/input/mean.binaryproto"
  }
  data_param {
    source: "/home/ML/cifar10/input/lmdb/valid_lmdb"
    batch_size: 50
    backend: LMDB
  }
}
```

-:--- train\_val\_3\_miniVggNet.prototxt 14% L67 (Fundamental)

```
bottom: "drop3"
top: "fc2"
param {
  lr_mult: 1
  decay_mult: 1
}
param {
  lr_mult: 2
  decay_mult: 0
}
inner_product_param {
  num_output: 10
  weight_filler {
    #type: "gaussian"
    type: "xavier"
    #std: 0.001
  }
  bias_filler {
    type: "constant"
    value: 1
  }
}
layer {
  name: "loss"
  type: "SoftmaxWithLoss"
  bottom: "fc2"
  bottom: "label"
  top: "loss"
}
layer {
  name: "accuracy"
  type: "Accuracy"
  bottom: "fc2"
  bottom: "label"
  top: "accuracy"
  include {
    phase: TEST
  }
}
```

-:--- train\_val\_3\_miniVggNet.prototxt 91% L519

# solver\_3\_miniVggNet.prototxt

```
net: "/home/ML/cifar10/caffe/models/miniVggNet/m3/train_val_3_miniVggNet.prototxt"
#
test_iter: 180          # test_iter = test dataset size / test batch size
#
test_interval: 1000     # amount of iterations after which the NN will test the performance on the test dataset
#
base_lr: 0.01           # the beginning of learning rate
#
lr_policy: "poly"       # it could be "step", "fixed", "exp", "poly", "sigmoid"
power: 1
#
#gamma: 0.1             # how much the learning rate should be changed every time we reach the next step
#
#stepsize: 8000
display: 100
max_iter: 40000         # end of NN training. Note that max_iter = num_epochs * training set size / test batch size
#
momentum: 0.9
weight_decay: 0.0005
snapshot: 10000
#snapshot_format: HDF5
snapshot_prefix: "/home/ML/cifar10/caffe/models/miniVggNet/m3/snapshot_3_miniVggNet_"
#
#solver_mode: CPU
solver_mode: GPU
#
#type: "Nesterov"
type: "SGD"

random_seed: 1201
```

-:--- solver\_3\_miniVggNet.prototxt All L1 (Fundamental)




# logfile\_3\_miniVggNet.log file (first fragment of training)

```
1 I0109 17:13:05.404975 13328 caffe.cpp:204] Using GPUs 0
2 I0109 17:13:05.430150 13328 caffe.cpp:209] GPU 0: Quadro P6000
3 I0109 17:13:05.703902 13328 solver.cpp:45] Initializing solver from parameters:
4 test_iter: 180
5 test_interval: 1000
6 base_lr: 0.01
7 display: 100
8 max_iter: 40000
9 lr_policy: "poly"
10 power: 1
11 momentum: 0.9
12 weight_decay: 0.0005
13 snapshot: 10000
14 snapshot_prefix: "/home/ML/cifar10/caffe/models/miniVggNet/m3/snapshot_3_miniVggNet_"
15 solver_mode: GPU
16 device_id: 0
17 random_seed: 1201
18 net: "/home/ML/cifar10/caffe/models/miniVggNet/m3/train_val_3_miniVggNet.prototxt"
19 train_state {
20   level: 0
21   stage: ""
22 }
23 type: "SGD"
24 I0109 17:13:05.704044 13328 solver.cpp:102] Creating training net from net file: /home/ML/cifar10/caffe/models/miniVggNet/m3/train_val_3_miniVggNet.prototxt
25 I0109 17:13:05.704278 13328 upgrade_proto.cpp:79] Attempting to upgrade batch norm layers using deprecated params: /home/ML/cifar10/caffe/models/miniVggNet/m3/train_val_3_miniVggNet.p
26 I0109 17:13:05.704286 13328 upgrade_proto.cpp:82] Successfully upgraded batch norm layers using deprecated params.
27 I0109 17:13:05.704350 13328 net.cpp:294] The NetState phase (0) differed from the phase (1) specified by a rule in layer data
28 I0109 17:13:05.704363 13328 net.cpp:294] The NetState phase (0) differed from the phase (1) specified by a rule in layer accuracy
29 I0109 17:13:05.704365 13328 net.cpp:294] The NetState phase (0) differed from the phase (1) specified by a rule in layer accuracy-top1
30 I0109 17:13:05.704368 13328 net.cpp:294] The NetState phase (0) differed from the phase (1) specified by a rule in layer accuracy-top5
31 I0109 17:13:05.704507 13328 net.cpp:51] Initializing net from parameters:
32 name: "miniVggNet on Cifar10 m3 NO-inPlace"
33 state {
34   phase: TRAIN
35   level: 0
36   stage: ""
37 }
38 layer {
39   name: "data"
40   type: "Data"
41   top: "data"
42   top: "label"
43   include {
```

# logfile\_3\_miniVggNet.log file (last fragment of training)

```
2957 I0109 17:25:31.738238 13328 sgd_solver.cpp:112] Iteration 39000, lr = 0.00025
2958 I0109 17:25:32.860574 13336 data_layer.cpp:73] Restarting data prefetching from start.
2959 I0109 17:25:33.639510 13328 solver.cpp:239] Iteration 39100 (52.5988 iter/s, 1.90119s/100 iters), loss = 0.173129
2960 I0109 17:25:33.639539 13328 solver.cpp:258] Train net output #0: loss = 0.173129 (* 1 = 0.173129 loss)
2961 I0109 17:25:33.639544 13328 sgd_solver.cpp:112] Iteration 39100, lr = 0.000225
2962 I0109 17:25:35.562705 13328 solver.cpp:239] Iteration 39200 (52.0001 iter/s, 1.92307s/100 iters), loss = 0.111885
2963 I0109 17:25:35.562731 13328 solver.cpp:258] Train net output #0: loss = 0.111885 (* 1 = 0.111885 loss)
2964 I0109 17:25:35.562736 13328 sgd_solver.cpp:112] Iteration 39200, lr = 0.0002
2965 I0109 17:25:37.426491 13328 solver.cpp:239] Iteration 39300 (53.6575 iter/s, 1.86367s/100 iters), loss = 0.136369
2966 I0109 17:25:37.426517 13328 solver.cpp:258] Train net output #0: loss = 0.136369 (* 1 = 0.136369 loss)
2967 I0109 17:25:37.426520 13328 sgd_solver.cpp:112] Iteration 39300, lr = 0.000175
2968 I0109 17:25:39.306041 13328 solver.cpp:239] Iteration 39400 (53.2075 iter/s, 1.87943s/100 iters), loss = 0.184061
2969 I0109 17:25:39.306092 13328 solver.cpp:258] Train net output #0: loss = 0.184061 (* 1 = 0.184061 loss)
2970 I0109 17:25:39.306097 13328 sgd_solver.cpp:112] Iteration 39400, lr = 0.00015
2971 I0109 17:25:40.229164 13336 data_layer.cpp:73] Restarting data prefetching from start.
2972 I0109 17:25:41.183041 13328 solver.cpp:239] Iteration 39500 (53.2805 iter/s, 1.87686s/100 iters), loss = 0.129273
2973 I0109 17:25:41.183068 13328 solver.cpp:258] Train net output #0: loss = 0.129273 (* 1 = 0.129273 loss)
2974 I0109 17:25:41.183071 13328 sgd_solver.cpp:112] Iteration 39500, lr = 0.000125
2975 I0109 17:25:43.079367 13328 solver.cpp:239] Iteration 39600 (52.7368 iter/s, 1.89621s/100 iters), loss = 0.118099
2976 I0109 17:25:43.079393 13328 solver.cpp:258] Train net output #0: loss = 0.118099 (* 1 = 0.118099 loss)
2977 I0109 17:25:43.079397 13328 sgd_solver.cpp:112] Iteration 39600, lr = 9.99999e-05
2978 I0109 17:25:44.965859 13328 solver.cpp:239] Iteration 39700 (53.0117 iter/s, 1.88638s/100 iters), loss = 0.116446
2979 I0109 17:25:44.965886 13328 solver.cpp:258] Train net output #0: loss = 0.116446 (* 1 = 0.116446 loss)
2980 I0109 17:25:44.965890 13328 sgd_solver.cpp:112] Iteration 39700, lr = 7.49999e-05
2981 I0109 17:25:46.821291 13328 solver.cpp:239] Iteration 39800 (53.8992 iter/s, 1.85532s/100 iters), loss = 0.125056
2982 I0109 17:25:46.821318 13328 solver.cpp:258] Train net output #0: loss = 0.125056 (* 1 = 0.125056 loss)
2983 I0109 17:25:46.821322 13328 sgd_solver.cpp:112] Iteration 39800, lr = 5e-05
2984 I0109 17:25:47.552762 13336 data_layer.cpp:73] Restarting data prefetching from start.
2985 I0109 17:25:48.683285 13328 solver.cpp:239] Iteration 39900 (53.7092 iter/s, 1.86188s/100 iters), loss = 0.177063
2986 I0109 17:25:48.683310 13328 solver.cpp:258] Train net output #0: loss = 0.177063 (* 1 = 0.177063 loss)
2987 I0109 17:25:48.683315 13328 sgd_solver.cpp:112] Iteration 39900, lr = 2.5e-05
2988 I0109 17:25:50.525246 13328 solver.cpp:464] Snapshotting to binary proto file /home/ML/cifar10/caffe/models/miniVggNet/m3/snapshot_3_miniVggNet_iter_40000.caffemodel
2989 I0109 17:25:50.549196 13328 sgd_solver.cpp:284] Snapshotting solver state to binary proto file /home/ML/cifar10/caffe/models/miniVggNet/m3/snapshot_3_miniVggNet_iter_40000.solverstate
2990 I0109 17:25:50.569916 13328 solver.cpp:327] Iteration 40000, loss = 0.133656
2991 I0109 17:25:50.569937 13328 solver.cpp:347] Iteration 40000, Testing net (#0)
2992 I0109 17:25:50.975668 13337 data_layer.cpp:73] Restarting data prefetching from start.
2993 I0109 17:25:50.983530 13328 solver.cpp:414] Test net output #0: accuracy = 0.857111
2994 I0109 17:25:50.983548 13328 solver.cpp:414] Test net output #1: loss = 0.449977 (* 1 = 0.449977 loss)
2995 I0109 17:25:50.983551 13328 solver.cpp:414] Test net output #2: top-1 = 0.857111
2996 I0109 17:25:50.983554 13328 solver.cpp:414] Test net output #3: top-5 = 0.991889
2997 I0109 17:25:50.983556 13328 solver.cpp:332] Optimization Done.
2998 I0109 17:25:50.983559 13328 caffe.cpp:250] Optimization Done.
2999
```

# miniVggNet: learning curves



# Average top1 accuracy on the test dataset: 87%

```
4633 IMAGE: /home/ML/cifar10/input/cifar10_jpg/test/deer_00076.jpg
4634 PREDICTED: 5
4635 EXPECTED : 4
4636 -----
4637 IMAGE: /home/ML/cifar10/input/cifar10_jpg/test/horse_00085.jpg
4638 PREDICTED: 7
4639 EXPECTED : 7
4640 -----
4641 IMAGE: /home/ML/cifar10/input/cifar10_jpg/test/ship_00003.jpg
4642 PREDICTED: 8
4643 EXPECTED : 8
4644 -----
4645 IMAGE: /home/ML/cifar10/input/cifar10_jpg/test/truck_00033.jpg
4646 PREDICTED: 9
4647 EXPECTED : 9
4648 -----
4649      precision    recall  f1-score   support
4650
4651     airplane      0.89      0.87      0.88        100
4652     automobile      0.91      0.94      0.93        100
4653         bird      0.83      0.82      0.82        100
4654         cat       0.73      0.70      0.71        100
4655         deer      0.81      0.83      0.82        100
4656         dog       0.83      0.83      0.83        100
4657         frog      0.87      0.90      0.88        100
4658         horse      0.99      0.94      0.96        100
4659         ship      0.89      0.93      0.91        100
4660         truck      0.94      0.92      0.93        100
4661
4662  avg / total      0.87      0.87      0.87       1000
4663
4664  SKLEARN Accuracy = 0.87
4665
4666
4667  TOP-5 ACCURACY              = 1.00
4668  TOP-5 FALSE                 = 0
4669  TOP-5 TRUE                   = 1000
4670
4671
4672  TOTAL NUMBER OF TRUE PREDICTIONS = 868
4673  TOTAL NUMBER OF FALSE PREDICTIONS = 132
4674  MANUALLY COMPUTED RECALL = 0.87
4675
```